

investigate-a-dataset-template-TMDB-Szymon-Debski

March 1, 2021

1 Project: Investigate the TMDB movie dataset.

1.1 Table of Contents

Introduction

Data Wrangling

Exploratory Data Analysis

Conclusions

Introduction

Dataset: TMDB movies

The dataset is based on 10,000 movies from The Movie Database (TMDb).

- In the analysis, we will be focusing on one hand on the runtime of the movies and their release year which we will be comparing to earnings and budgets respectively.

Questions:

- Does the runtime of movies impact their earnings?
- Do newer movies have bigger budgets?

```
[1]: import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import matplotlib.cm as cm
import seaborn as sns
import warnings
warnings.simplefilter(action='ignore', category=FutureWarning)

%matplotlib inline
```

Data Wrangling

In this section, I will load in the data, clean it and remove unnecessary columns and rows.

1.1.1 General Properties

Dataset I chose for this analysis: TMDb movies. In my analysis, I will focus on top-earning movies and movies with big budgets. I will try to answer which characteristics correspond to a high-earning and big budgets.

```
[39]: #data set loaded
```

```
df = pd.read_csv('tmdb-movies.csv')
df.head(1)
```

```
[39]:      id  imdb_id  popularity    budget    revenue  original_title \
0  135397  tt0369610    32.98576  150000000  1513528810  Jurassic World

                                     cast \
0  Chris Pratt|Bryce Dallas Howard|Irrfan Khan|Vi...

                                     homepage    director    tagline ... \
0  http://www.jurassicworld.com/  Colin Trevorrow  The park is open. ...

                                     overview runtime \
0  Twenty-two years after the events of Jurassic ...    124

                                     genres \
0  Action|Adventure|Science Fiction|Thriller

                                     production_companies  release_date  vote_count \
0  Universal Studios|Amblin Entertainment|Legenda...    6/9/15    5562

    vote_average  release_year    budget_adj    revenue_adj
0      6.50000      2015  137999939.28003  1392445892.52380

[1 rows x 21 columns]
```

1.1.2 Data Cleaning

- Dropped unneeded columns ('cast', 'homepage', 'tagline', 'keywords', 'overview', 'production_companies')
- Next I discarded missing values
- After that, I set the right date format
- Also I changed the number format which made the numbers more readable
- Next I cleaned the genres column so that it shows only the first genre
- Next step was to clean the duplicates (there was one)
- After that I deleted rows with 0 for 'budget_adj' and 'revenue_adj'. I reasoned that replacing the values with a mean would distort the data

1.2 In the end I was left with a data frame with 3853 rows and 15 columns

```
[40]: df.describe()
```

```
[40]:
```

	id	popularity	budget	revenue	runtime \
count	10866.00000	10866.00000	10866.00000	10866.00000	10866.00000
mean	66064.17743	0.64644	14625701.09415	39823319.79339	102.07086
std	92130.13656	1.00018	30913213.83144	117003486.58209	31.38141
min	5.00000	0.00006	0.00000	0.00000	0.00000
25%	10596.25000	0.20758	0.00000	0.00000	90.00000
50%	20669.00000	0.38386	0.00000	0.00000	99.00000
75%	75610.00000	0.71382	15000000.00000	24000000.00000	111.00000
max	417859.00000	32.98576	425000000.00000	2781505847.00000	900.00000

	vote_count	vote_average	release_year	budget_adj	revenue_adj
count	10866.00000	10866.00000	10866.00000	10866.00000	10866.00000
mean	217.38975	5.97492	2001.32266	17551039.82289	51364363.25325
std	575.61906	0.93514	12.81294	34306155.72284	144632485.03997
min	10.00000	1.50000	1960.00000	0.00000	0.00000
25%	17.00000	5.40000	1995.00000	0.00000	0.00000
50%	38.00000	6.00000	2006.00000	0.00000	0.00000
75%	145.75000	6.60000	2011.00000	20853251.08440	33697095.71731
max	9767.00000	9.20000	2015.00000	425000000.00000	2827123750.41189

```
[41]: df.shape
```

```
[41]: (10866, 21)
```

```
[42]: df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 10866 entries, 0 to 10865
Data columns (total 21 columns):
#   Column                Non-Null Count  Dtype
---  -
0   id                     10866 non-null  int64
1   imdb_id                10856 non-null  object
2   popularity             10866 non-null  float64
3   budget                 10866 non-null  int64
4   revenue                10866 non-null  int64
5   original_title         10866 non-null  object
6   cast                   10790 non-null  object
7   homepage               2936 non-null  object
8   director               10822 non-null  object
9   tagline                8042 non-null  object
10  keywords               9373 non-null  object
11  overview               10862 non-null  object
12  runtime                10866 non-null  int64
13  genres                 10843 non-null  object
```

```

14 production_companies 9836 non-null object
15 release_date          10866 non-null object
16 vote_count            10866 non-null int64
17 vote_average          10866 non-null float64
18 release_year          10866 non-null int64
19 budget_adj            10866 non-null float64
20 revenue_adj           10866 non-null float64
dtypes: float64(4), int64(6), object(11)
memory usage: 1.7+ MB

```

```

[43]: #dropped unused columns
df.drop(['cast', 'homepage', 'tagline', 'keywords', 'overview',
        ↪ 'production_companies'], axis=1, inplace=True)

```

```

[44]: df.head(1)

```

```

[44]:      id  imdb_id  popularity    budget    revenue  original_title \
0  135397  tt0369610    32.98576  150000000  1513528810  Jurassic World

      director  runtime                                genres \
0  Colin Trevorrow    124  Action|Adventure|Science Fiction|Thriller

      release_date  vote_count  vote_average  release_year    budget_adj \
0      6/9/15          5562        6.50000         2015  137999939.28003

      revenue_adj
0  1392445892.52380

```

```

[45]: # dropped missing values
df.dropna(inplace=True)

```

```

[46]: # changed date format
df['release_date'] = pd.to_datetime(df['release_date'], format='%m/%d/%y')

```

```

[47]: # changed number format
pd.set_option('display.float_format', lambda x: '%.5f' % x)

```

```

[48]: df.info()

```

```

<class 'pandas.core.frame.DataFrame'>
Int64Index: 10796 entries, 0 to 10865
Data columns (total 15 columns):
#   Column          Non-Null Count  Dtype
---  -
0   id              10796 non-null  int64
1   imdb_id         10796 non-null  object
2   popularity      10796 non-null  float64
3   budget          10796 non-null  int64

```

```

4  revenue          10796 non-null  int64
5  original_title   10796 non-null  object
6  director         10796 non-null  object
7  runtime          10796 non-null  int64
8  genres           10796 non-null  object
9  release_date     10796 non-null  datetime64[ns]
10 vote_count       10796 non-null  int64
11 vote_average     10796 non-null  float64
12 release_year     10796 non-null  int64
13 budget_adj       10796 non-null  float64
14 revenue_adj      10796 non-null  float64
dtypes: datetime64[ns](1), float64(4), int64(6), object(4)
memory usage: 1.3+ MB

```

```
[49]: def split(column):
      x = column.split('|')[0]
      return x
```

```
[52]: df['genres'] = df['genres'].apply(lambda x: split(x))
```

```
[53]: df.head(1)
```

```
[53]:
```

	id	imdb_id	popularity	budget	revenue	original_title	\
0	135397	tt0369610	32.98576	150000000	1513528810	Jurassic World	

	director	runtime	genres	release_date	vote_count	vote_average	\
0	Colin Trevorrow	124	Action	2015-06-09	5562	6.50000	

	release_year	budget_adj	revenue_adj
0	2015	137999939.28003	1392445892.52380

```
[54]: df.describe()
```

```
[54]:
```

	id	popularity	budget	revenue	runtime	\
count	10796.00000	10796.00000	10796.00000	10796.00000	10796.00000	
mean	65558.31808	0.64961	14719366.67182	40080510.64052	102.21332	
std	91747.96902	1.00258	30991238.22095	117338430.76150	30.76277	
min	5.00000	0.00019	0.00000	0.00000	0.00000	
25%	10568.50000	0.20920	0.00000	0.00000	90.00000	
50%	20454.00000	0.38551	0.00000	0.00000	99.00000	
75%	74663.50000	0.71772	16000000.00000	24609991.25000	112.00000	
max	417859.00000	32.98576	425000000.00000	2781505847.00000	900.00000	

	vote_count	vote_average	release_year	budget_adj	revenue_adj
count	10796.00000	10796.00000	10796.00000	10796.00000	10796.00000
mean	218.68164	5.97030	2001.28677	17663691.65299	51696371.84669
std	577.25738	0.93292	12.82103	34388506.64822	145041642.43934
min	10.00000	1.50000	1960.00000	0.00000	0.00000

25%	17.00000	5.40000	1995.00000	0.00000	0.00000
50%	39.00000	6.00000	2006.00000	0.00000	0.00000
75%	147.00000	6.60000	2011.00000	21033371.65263	34097666.53813
max	9767.00000	9.20000	2015.00000	425000000.00000	2827123750.41189

```
[55]: df.query('budget_adj == 0').info()
```

```
<class 'pandas.core.frame.DataFrame'>
Int64Index: 5632 entries, 30 to 10864
Data columns (total 15 columns):
#   Column                Non-Null Count  Dtype
---  -
0   id                    5632 non-null   int64
1   imdb_id              5632 non-null   object
2   popularity            5632 non-null   float64
3   budget               5632 non-null   int64
4   revenue              5632 non-null   int64
5   original_title       5632 non-null   object
6   director             5632 non-null   object
7   runtime              5632 non-null   int64
8   genres               5632 non-null   object
9   release_date         5632 non-null   datetime64[ns]
10  vote_count           5632 non-null   int64
11  vote_average         5632 non-null   float64
12  release_year         5632 non-null   int64
13  budget_adj           5632 non-null   float64
14  revenue_adj          5632 non-null   float64
dtypes: datetime64[ns](1), float64(4), int64(6), object(4)
memory usage: 704.0+ KB
```

```
[56]: # duplicates
sum(df.duplicated())
```

```
[56]: 1
```

```
[57]: df.shape
```

```
[57]: (10796, 15)
```

```
[58]: # dropped duplicates
df.drop_duplicates(keep='first', inplace=True)
```

```
[59]: df.shape
```

```
[59]: (10795, 15)
```

```
[60]: # deleted rows with 0 in 'budget_adj' column
df = df[df['budget_adj'] != 0]
```

```
[61]: df.shape
```

```
[61]: (5163, 15)
```

```
[62]: # deleted rows with 0 in 'revenue_adj' column  
df = df[df['revenue_adj'] != 0]
```

```
[63]: df.shape
```

```
[63]: (3853, 15)
```

Exploratory Data Analysis

Now that my data is clean I will research my questions.

1.3 Does the runtime of movies impact their earnings?

```
[64]: df['earnings'] = df['revenue_adj'] - df['budget_adj']
```

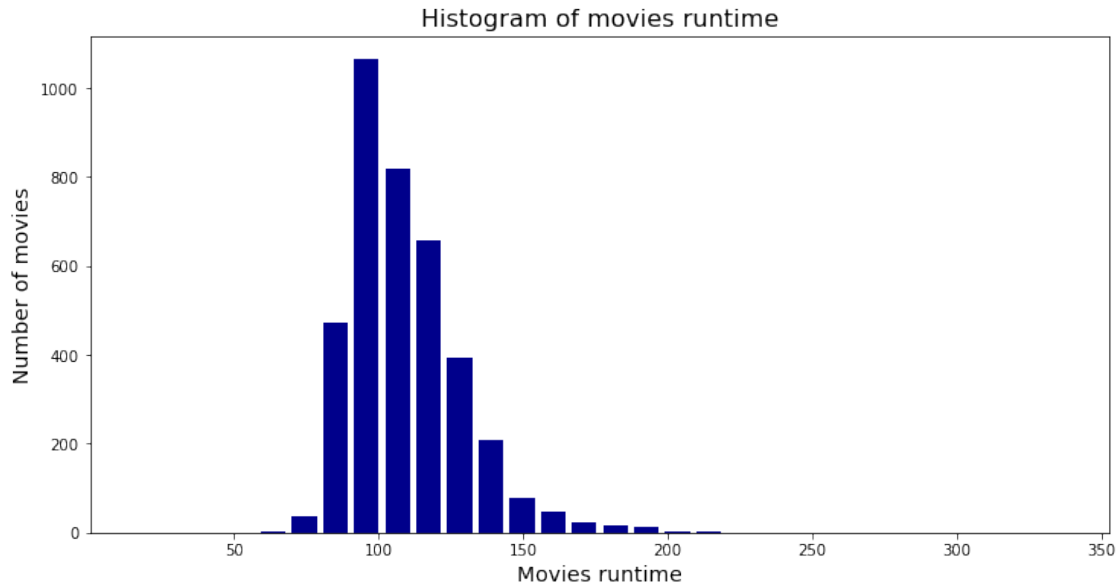
```
[65]: df.head(1)
```

```
[65]:      id  imdb_id  popularity    budget    revenue  original_title \  
0  135397  tt0369610    32.98576  150000000  1513528810  Jurassic World  
  
      director  runtime  genres  release_date  vote_count  vote_average \  
0  Colin Trevorrow    124  Action   2015-06-09         5562         6.50000  
  
      release_year    budget_adj    revenue_adj    earnings  
0         2015  137999939.28003  1392445892.52380  1254445953.24377
```

```
[66]: df.shape
```

```
[66]: (3853, 16)
```

```
[67]: plt.figure(figsize=(12,6))  
  
plt.xlabel('Movies runtime', fontsize=14)  
plt.ylabel('Number of movies', fontsize=14)  
plt.title('Histogram of movies runtime', fontsize=16)  
plt.hist(df['runtime'], rwidth = 0.8, bins=30, color='darkblue')  
plt.show()
```



We can see that the runtime distribution is skewed to the right. Most of the movies have a runtime of 90 - 100 minutes.

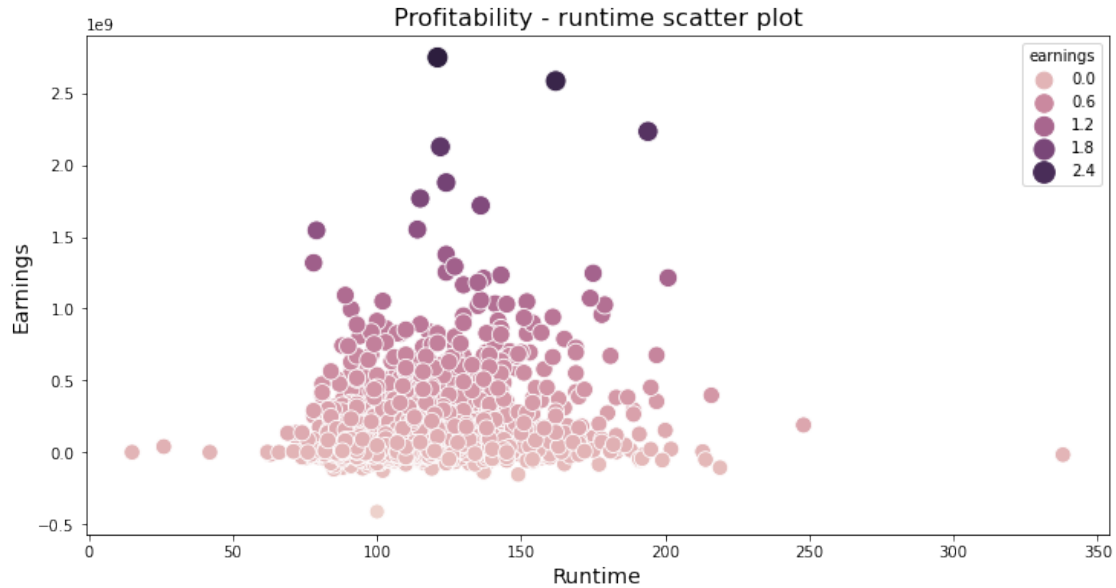
```
[68]: df.runtime.describe()
```

```
[68]: count    3853.00000
      mean      109.20893
      std       19.91291
      min       15.00000
      25%       95.00000
      50%      106.00000
      75%      119.00000
      max      338.00000
      Name: runtime, dtype: float64
```

```
[69]: plt.figure(figsize=(12,6))

      plt.xlabel('Runtime', fontsize=14)
      plt.ylabel('Earnings', fontsize=14)
      plt.title('Profitability - runtime scatter plot', fontsize=16)

      sns.scatterplot('runtime', 'earnings', data=df, hue='earnings',
                      size='earnings', sizes=(100, 200))
      plt.show()
```

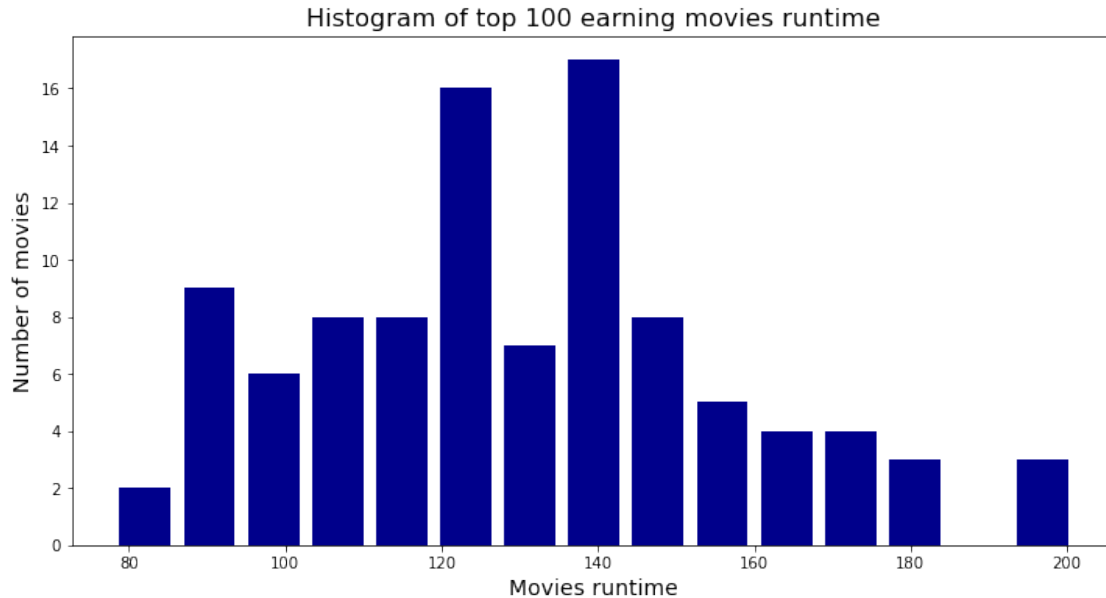



Above we can see a scatter plot of runtime and earnings. From the graph above we can see that the most profitable movies have a runtime greater than the mean runtime. Let's analyze it a little bit more by selecting only the top-earning movies.

```
[70]: # created a new df only for the 100 to earners
top_100_earnings = df.sort_values(by=['earnings'], ascending = False).head(100)
```

```
[71]: plt.figure(figsize=(12,6))

plt.xlabel('Movies runtime', fontsize=14)
plt.ylabel('Number of movies', fontsize=14)
plt.title('Histogram of top 100 earning movies runtime', fontsize=16)
plt.hist(top_100_earnings['runtime'], rwidth = 0.8, bins=15, color='darkblue')
plt.show()
```



We can see from the graph above that the runtime distribution is different for top earners. It is still skewed to the right but not so much so. The biggest difference is that in this histogram most of the top-earning movies have a runtime of 140 min compared to 100 for all analyzed movies. The runtime mean for top earners is 130 minutes vs 109 for all analyzed movies.

```
[72]: top_100_earnings.runtime.describe()
```

```
[72]: count    100.00000
      mean     130.50000
      std      26.66231
      min      78.00000
      25%     113.00000
      50%     129.50000
      75%     145.25000
      max     201.00000
      Name: runtime, dtype: float64
```

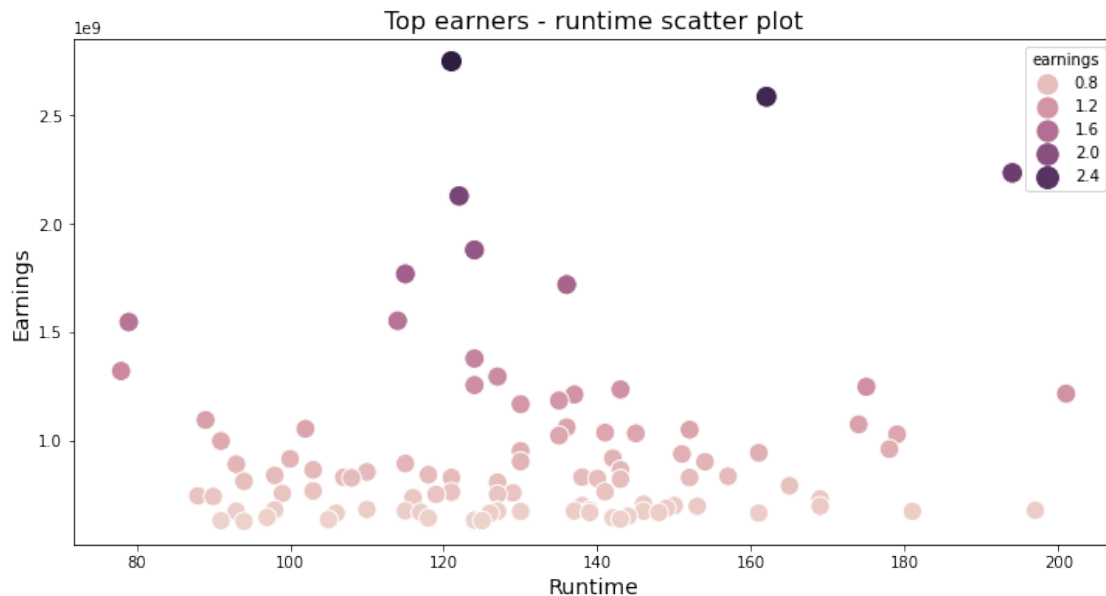
```
[73]: plt.figure(figsize=(12,6))

      plt.xlabel('Runtime', fontsize=14)
      plt.ylabel('Earnings', fontsize=14)
      plt.title('Top earners - runtime scatter plot', fontsize=16)

      sns.scatterplot('runtime', 'earnings', data=top_100_earnings, hue='earnings',
                      size='earnings', sizes=(150, 200))

      # plt.scatter(top_100_earnings['runtime'], top_100_earnings['earnings'])
```

```
plt.show()
```



When looking at the graph above we can see that most of the top-earning movies have a runtime greater than 130 minutes.

1.3.1 We can say that although runtime does not determine that a movie will be profitable (there are movies that are long and did not earn too much or even had a loss), we can conclude that longer movies have a better chance of being profitable. This may be connected with the fact that blockbusters are generally long and have big budgets.

1.4 Do newer movies have bigger budgets?

```
[74]: df.head(1)
```

```
[74]:      id  imdb_id  popularity    budget    revenue  original_title \
0  135397  tt0369610    32.98576  150000000  1513528810  Jurassic World

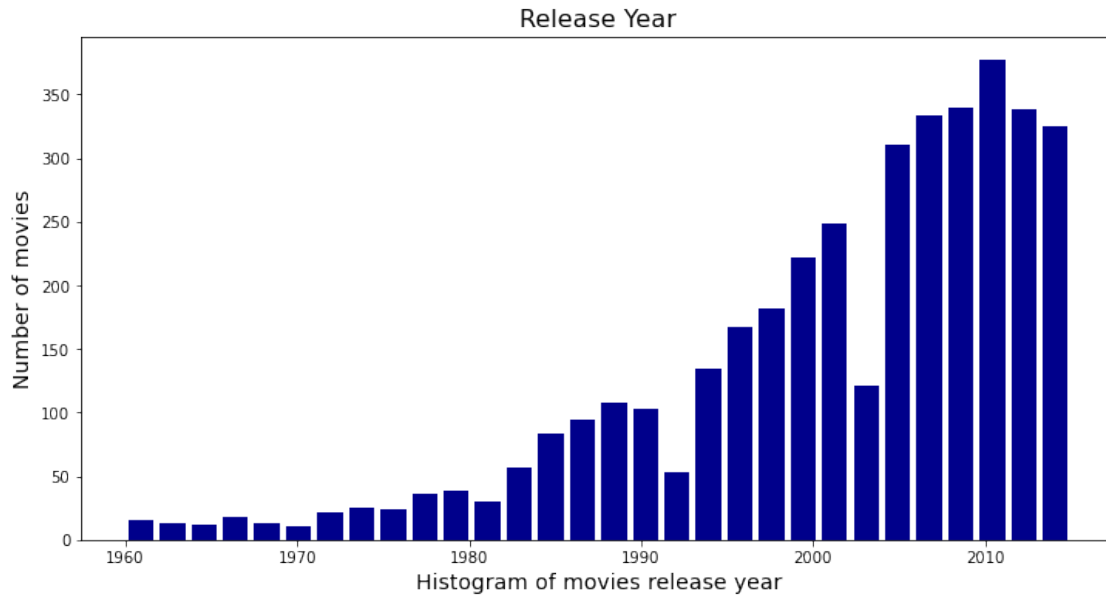
      director  runtime  genres  release_date  vote_count  vote_average \
0  Colin Trevorrow    124  Action    2015-06-09         5562         6.50000

      release_year    budget_adj    revenue_adj    earnings
0         2015  137999939.28003  1392445892.52380  1254445953.24377
```

```
[75]: plt.figure(figsize=(12,6))

plt.xlabel('Histogram of movies release year', fontsize=14)
```

```
plt.ylabel('Number of movies', fontsize=14)
plt.title('Release Year', fontsize=16)
plt.hist(df['release_year'], rwidth = 0.8, bins=30, color='darkblue')
plt.show()
```



We can see that the distribution is heavily skewed to the left which is understandable - more and more movies are released every year.

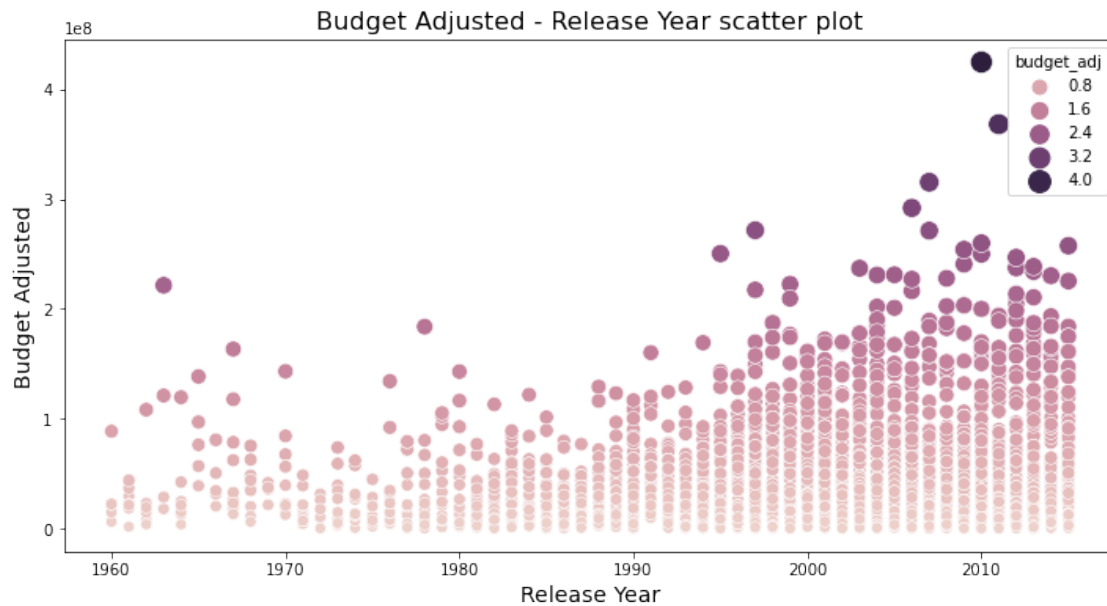
```
[76]: df.release_year.describe()
```

```
[76]: count    3853.00000
      mean    2001.25928
      std      11.28352
      min    1960.00000
      25%    1995.00000
      50%    2004.00000
      75%    2010.00000
      max    2015.00000
      Name: release_year, dtype: float64
```

```
[77]: plt.figure(figsize=(12,6))

plt.xlabel('Release Year', fontsize=14)
plt.ylabel('Budget Adjusted', fontsize=14)
plt.title('Budget Adjusted - Release Year scatter plot', fontsize=16)
# plt.scatter(df['release_year'], df['earnings'])
sns.scatterplot('release_year', 'budget_adj', data=df, hue='budget_adj',
                size='budget_adj', sizes=(50, 200))
```

```
plt.show()
```



Above we can see a scatter plot of release year and budget adjusted. There is a clear indication that movies released after 1995 have a bigger budget. We can see that the movies that had the biggest budgets were released after 1995. It is worth mentioning that I'm using the budget adjusted values which means inflation did not impact the results. ## Conclusions

To answer the first question we can say that: longer movies have a better chance of being profitable.

Answering the second question we can definitely say the newer movies have bigger budgets.

1.5 Limitations:

1. We base most of our analysis on the budget_adj and revenue_adj columns. Unfortunately we don't know how they were adjusted - what were the exact assumptions. If we knew the exact calculation we could have adjusted our analysis accordingly.
2. Due to many rows of missing data for budget and revenue, we had to drop around 7000 rows which are not very good for our analysis. This limitation decreased the accuracy of the analysis tremendously.
3. The last limitation in my opinion is that data is not updated. In recent years a lot has changed in the movie industry. The newest movies in the data set are from 2015.

[]: